

<b>Started on</b>	Thursday, 12 June 2025, 4:53 PM
<b>State</b>	Finished
<b>Completed on</b>	Thursday, 12 June 2025, 4:58 PM
<b>Time taken</b>	4 mins 44 secs
<b>Marks</b>	8.00/15.00
<b>Grade</b>	<b>53.33</b> out of 100.00

**Question 1**

Complete

Mark 1.00 out of 1.00

In a standard RNN, the hidden state  $h_t$  is updated as:

- ☐ a.  $h_t = \text{ReLU}(x_t)$
- ☐ b.  $h_t = \sigma(W x_t + b)$
- ☒ c.  $h_t = \tanh(W x_t + U h_{t-1} + b)$
- ☐ d.  $h_t = \tanh(W x_t + b)$

**Question 2**

Complete

Mark 1.00 out of 1.00

In an LSTM cell, what is the function of the cell state  $C_t$ ?

- ☐ a. Acts as the output layer
- ☐ b. Stores hidden layers
- ☐ c. Calculates gradients
- ☒ d. Stores long-term memory

**Question 3**

Complete

Mark 0.00 out of 1.00

In an LSTM cell, which gate controls how much of the previous hidden state should be carried forward?

- ☐ a. Forget gate
- ☐ b. Input gate
- ☒ c. Memory gate
- ☐ d. Output gate

**Question 4**

Complete

Mark 1.00 out of 1.00

In sequence-to-sequence models, what is the role of the encoder?

- ☒ a. Encode input sequence into a fixed representation
- ☐ b. Predict next token
- ☐ c. Update output vocabulary
- ☐ d. Translate output sequence

**Question 5**

Complete

Mark 1.00 out of 1.00

What does teacher forcing refer to during RNN training?

- ☐ a. Resetting hidden states between batches
- ☒ b. Feeding the ground truth output at time  $t-1$  to predict time  $t$
- ☐ c. Using the model's own output as input
- ☐ d. Pre-training the encoder before decoder

**Question 6**

Complete

Mark 1.00 out of 1.00

What is gradient clipping in the context of training RNNs?

- ☐ a. Limiting updates to only the final layer
- ☒ b. Restricting the magnitude of gradients to prevent exploding gradients
- ☐ c. Reducing batch size to avoid overfitting
- ☐ d. Applying dropout to avoid vanishing gradients

**Question 7**

Complete

Mark 0.00 out of 1.00

What is the main reason RNNs struggle with learning long-term dependencies?

- ☐ a. Insufficient parameters
- ☒ b. Gradient explosion
- ☐ c. Vanishing gradients
- ☐ d. Lack of activation functions

**Question 8**

Complete

Mark 1.00 out of 1.00

What is the primary advantage of using bidirectional RNNs?

- ☒ a. Access to both past and future context
- ☐ b. Replaces the need for attention mechanisms
- ☐ c. Works with images
- ☐ d. Reduced computation time

**Question 9**

Complete

Mark 0.00 out of 1.00

What technique is commonly used during inference in seq2seq models to improve generation quality?

- ☐ a. Adam optimizer
- ☒ b. Batch normalization
- ☐ c. Beam search
- ☐ d. Dropout

**Question 10**

Complete

Mark 0.00 out of 1.00

Which loss function is most commonly used in training sequence-to-sequence models with RNNs for classification?

- ☐ a. Binary Crossentropy
- ☐ b. Categorical Crossentropy
- ☐ c. Mean Squared Error
- ☒ d. Hinge Loss

**Question 11**

Complete

Mark 0.00 out of 1.00

Which mechanism allows RNN-based models to focus on specific parts of the input during decoding?

- ☐ a. Attention
- ☐ b. Dropout
- ☒ c. Beam search
- ☐ d. Batch normalization

**Question 12**

Complete

Mark 0.00 out of 1.00

Which of the following statements about GRU is incorrect?

- ☐ a. GRU is generally faster to train than LSTM
- ☐ b. GRU has fewer parameters than LSTM
- ☒ c. GRU combines the forget and input gates into a single update gate
- ☐ d. GRU has a separate memory cell  $c_t$  like LSTM

**Question 13**

Complete

Mark 0.00 out of 1.00

Which one is not a typical application of RNNs?

- ☐ a. Object detection
- ☐ b. Machine translation
- ☒ c. Sentiment analysis
- ☐ d. Speech recognition

**Question 14**

Complete

Mark 1.00 out of 1.00

Which RNN variant is specifically designed to solve the vanishing gradient problem?

- ☐ a. Vanilla RNN
- ☒ b. LSTM
- ☐ c. Bidirectional RNN
- ☐ d. GRU

**Question 15**

Complete

Mark 1.00 out of 1.00

Why are RNNs not inherently parallelizable across time steps?

- ☐ a. They use convolutional filters
- ☒ b. Each output depends on previous output
- ☐ c. Due to weight sharing
- ☐ d. They have attention layers

